

MORTIMER: Multi-asset Optimisation and Risk Targeting with Integrated Monthly E6 Rebalancing

Nelson Alvarez, Amir Hajomar, Sean Majewski, and Miles Rack

4 October 2026

Summary

MORTIMER combines monthly equal-risk-contribution (ERC) weights with a causal volatility overlay across six exchange-traded funds (ETFs). In the held-out out-of-sample (OOS) test, the portfolio returned 9.52% annualised after five basis points of trading costs, with 7.47% volatility, a 0.714 Sharpe ratio, a -7.00% maximum drawdown, and annual half-L1 turnover of 1.119. Relative to matched unscaled ERC, volatility fell by 8.2% and drawdown by 9.9%, while the Sharpe ratio also declined. Validation showed larger risk reductions alongside a 15-basis-point annual return difference. The results indicate that volatility targeting moderated portfolio risk without improving risk-adjusted return.

Economic hypothesis

Investors accept equity, credit, duration, and commodity risk premia, while benchmark mandates and slow risk-budget adjustment can leave portfolios exposed as volatility rises. Since variance clusters, a forecast based on recent returns may help manage subsequent portfolio risk without predicting return direction. MORTIMER reduces risky exposure when forecast volatility exceeds its target. We test whether this rule lowers volatility and tail losses after costs, while limiting the associated return drag. A shock can arrive before the overlay responds, and lower exposure can delay recovery.

Data and universe

The fixed universe comprises QQQ, IWM, HYG, TLT, GLD, and DBC. Yahoo adjusted closes include distributions and splits; raw closes and volume support liquidity estimates. Cash accrues from initial-release observations in the Archival Federal Reserve Economic Data (ALFRED) DGS3MO series, using the next exchange session for availability and actual/365 day-count accrual. The New York Stock Exchange (NYSE) calendar spans 2 January 2008 to 2 October 2026, with 550 sessions reserved for estimation before scored returns. Training ends on 31 December 2018, and validation runs from 1 January 2019 to 1 October 2024. The held-out OOS period runs from 2 October 2024 to 2 October 2026, the shorter of the most recent 20% of sessions and two years. The six funds were selected for continuous price and volume histories through the end of the sample; membership is fixed throughout the analysis.

Methodology

At each month-end, estimate daily covariance from 126 completed observations using Ledoit–Wolf shrinkage toward a scaled-identity target (Ledoit and Wolf, 2004). The constrained equal-risk-contribution (ERC) allocation for six assets solves

$$\min_b \sum_{i=1}^6 \left[b_i (\hat{\Sigma}_t b)_i - \frac{b^\top \hat{\Sigma}_t b}{6} \right]^2, \quad \mathbf{1}^\top b = 1, \quad 0 \leq b_i \leq 0.35.$$

Here b_i is the base weight, $\hat{\Sigma}_t$ is daily return covariance, and $b_i (\hat{\Sigma}_t b)_i$ is variance contribution. Binding name caps can prevent exact equality.

For current base b_t , calculate historical portfolio returns $x_s = b_t^\top r_s$ using observations through close t . An exponentially weighted moving average (EWMA) with span n applies geometrically declining weights to squared returns. The annualised volatility forecast and capped risky exposure are

$$\hat{\sigma}_t = \sqrt{252 \max\{\text{EWMA}_{20}(x_s^2), \text{EWMA}_{63}(x_s^2)\}}, \quad k_t = \min(1, 0.10/\hat{\sigma}_t).$$

Each estimate requires 63 observations. Taking the larger estimate combines a faster response with a slower volatility measure. A five-percentage-point exposure change triggers a trade, while monthly base rebalances proceed independently. Gross exposure cannot exceed one, and the strategy neither shorts nor borrows. The 10% forecast target does not fix realised volatility.

The close- t signal is executed at close $t + 1$, after that session’s return; new holdings first earn the next close-to-close return. Given pre-fee risky holdings h_i , portfolio net asset value (NAV) V , desired weights w_i , and cost rate c , post-fee NAV satisfies


$$V^{post} = V - c \sum_i |V^{post} w_i - h_i|.$$

Costs are five basis points of risky dollars bought plus sold, representing spread, commission, and closing slippage; cash transfers carry no charge. A ten-basis-point cost case is evaluated on development data. The cash return proxy is derived from initial-release three-month Treasury bill yields.

Results

We calculate annualised return as compound annual growth rate (CAGR) from split-local net wealth using 252 trading sessions per year. Annualised volatility is the sample standard deviation of daily returns multiplied by $\sqrt{252}$, and the Sharpe ratio uses daily excess returns relative to the cash return proxy. Maximum drawdown is measured from initial wealth of one. Annual turnover is half the L1 change in risky and cash weights, and transaction costs apply to risky purchases and sales. The table reports each metric net of five-basis-point transaction costs for training, validation, and the held-out out-of-sample test; the adjacent panels show net equity curves and drawdowns for validation and test.


Sample	Rule	Ann. return%	Vol.%	Sharpe	Sortino	MDD%	Calmar	Turnover
Training	MORTIMER	5.85	6.37	0.859	1.232	-12.36	0.473	0.96
Training	Unscaled ERC	6.13	6.51	0.884	1.273	-12.48	0.491	0.45
Validation	MORTIMER	7.56	7.93	0.660	0.916	-16.63	0.455	1.57
Validation	Unscaled ERC	7.71	9.78	0.566	0.781	-19.14	0.403	0.35
OOS test	MORTIMER	9.52	7.47	0.714	0.984	-7.00	1.360	1.12
OOS test	Unscaled ERC	11.00	8.13	0.827	1.181	-7.77	1.416	0.35



The volatility overlay has its clearest effect in validation. Relative to unscaled ERC, annualised volatility falls from 9.78% to 7.93%, and maximum drawdown from -19.14% to -16.63%; annualised return is lower by 15 basis points. At ten basis points per trade, validation CAGR and Sharpe are 7.46% and 0.649. A paired 20-session block bootstrap estimates a 95% interval of [-2.94, 2.44] percentage points for the validation return difference. In the held-out OOS test, volatility and maximum drawdown are lower by 8.2% and 9.9%, respectively; the Sharpe ratio is 0.714, compared with 0.827 for unscaled ERC. The forecast's validation Pearson and rank information coefficients are 0.537 and 0.668 against following-window realised volatility.

Risk management

Base weights are long-only and capped at 35% per instrument. Gross exposure cannot exceed 100%, with the remainder held as cash. Weights drift between executions and may exceed target caps. The overlay reduces or restores exposure when the forecast changes by five percentage points. Ordinary least squares (OLS) attribution on 479 common sessions estimates betas of 0.272 to the market, 0.031 to value, 0.066 to momentum, and 0.198 to Treasury-duration returns.



Liquidity and capacity

Liquidity is measured with previous-session trailing 20-session average daily dollar volume (ADDV). Participation is Q/ADDV , where Q is risky trade notional. The square-root impact scenario is

$$I = Y\sigma\sqrt{Q/\text{ADDV}}, \quad \text{NAV cost} = I Q/\text{NAV},$$

We set $Y = 0.5$ and use lagged volatility σ for each AUM scenario.

Initial AUM (\$m)	Max $Q/\text{ADDV}\%$	Impact-net CAGR%
1	0.52	7.55
10	5.22	7.53
100	52.16	7.47
500	260.81	7.35
1000	521.62	7.26

At $Y = 0.5$, a 1% participation limit is reached at initial AUM of \$1.92 million. The higher AUM rows exceed 100% participation and show the impact formula extrapolated beyond that level.

Limitations and further work

The portfolio history uses funds that survive to the end of the sample, so it does not capture discontinued funds. Further work can extend the analysis with point-in-time fund histories and transaction-level execution data, and test the capacity estimates against observed market impact and extend the doubled-cost sensitivity using empirically calibrated spreads and fees.

References

- Brixton, A., Brooks, J., Hecht, P., Ilmanen, A., Maloney, T., and McQuinn, N. (2023). *A changing stock–bond correlation: drivers and implications*. AQR Capital Management. <https://www.aqr.com/insights/research/journal-article/a-changing-stock-bond-correlation>. Accessed 4 October 2026.
- French, K. R. (n.d.). *Data library*. Dartmouth College. https://mba.tuck.dartmouth.edu/pages/faculty/ken.french/data_library.html. Accessed 4 October 2026.
- Federal Reserve Bank of St. Louis (n.d.). *Series observations*. FRED and ALFRED documentation. https://fred.stlouisfed.org/docs/api/fred/series_observations.html. Accessed 4 October 2026.
- Federal Reserve Bank of St. Louis (n.d.). *DGS3MO: market yield on US Treasury securities at three-month constant maturity*. FRED. <https://fred.stlouisfed.org/series/DGS3MO>. Accessed 4 October 2026.
- Ledoit, O. and Wolf, M. (2004). *A well-conditioned estimator for large-dimensional covariance matrices*. Journal of Multivariate Analysis, 88(2), 365–411. [https://doi.org/10.1016/S0047-259X\(03\)00096-4](https://doi.org/10.1016/S0047-259X(03)00096-4). Accessed 4 October 2026.
- Moreira, A. and Muir, T. (2017). *Volatility-managed portfolios*. Journal of Finance, 72(4), 1611–1644. <https://doi.org/10.1111/jofi.12513>. Accessed 4 October 2026.
- Yahoo Finance (n.d.). *Historical market prices, adjustments, and volume*. Yahoo Finance. <https://finance.yahoo.com/>. Accessed 4 October 2026.